

# Assignment 3 : Predict 454

*Rahul Sangole*

*2018-11-05*

```
library('tidyverse')
library('knitr')
library('kableExtra')
library('foreach')
library("doMC")
library("doParallel")
library('ggthemr')
ggthemr("fresh")
options(knitr.table.format = "latex")
knitr::opts_chunk$set(cache = TRUE)
```

## Loading the data

Loading the data into the system for the first time is attempted using two techniques. In the first, `purrr::map_df` (tidy equivalent of the `lapply`/`sapply` core functions) is used to iterate over the unzipped csv files in the folder, apply column types, merge into a tibble, and store in `df`. We can see that `df` is a large tibble, of 11.5 million observations by 30 columns. Though, this method results in the entire tibble being in memory.

```
df <- list.files(pattern = "(csv)$") %>%
  map_df(~read_csv(.x,
    col_types = cols(
      .default = col_integer(),
      UniqueCarrier = col_character(),
      TailNum = col_character(),
      AirTime = col_character(),
      Origin = col_character(),
      Dest = col_character(),
      TaxiIn = col_character(),
      TaxiOut = col_character(),
      CancellationCode = col_character(),
      CarrierDelay = col_character(),
      WeatherDelay = col_character(),
      NASDelay = col_character(),
      SecurityDelay = col_character(),
      LateAircraftDelay = col_character()
    ),
    col_names = T,
    progress = T))
```

```
##
```

```
| | 0%
| | 0%
| | 0%
| | 0%
| | 0%
```

```
dim(df)
```

```

    dbWriteTable(
      conn = con,
      name = 'airlines',
      value = df,
      append = TRUE,
      row.names = FALSE
    )
  }
  # close the connection
  dbDisconnect(con)

```

The third method - usage of `bigmemory` was attempted. Though, due to some error on my system, I could not get the read function to work correctly. Thus, I could not investigate it's usage.

## Comparison of %do and %dopar

Since I only have a four core machine, there's a chance that a plain %do might fare better than a %dopar given that the latter has some overhead associated with creation of the parallel executors. I'm going to test this first using `microbenchmark`. The calculations are so quick, that we must run  $10^6$  iterations to see some differences in execution time.

```

df$depHours <- floor(df$CRSDepTime/100)
df$depHours[df$depHours==24] <- 0
splits <- split(1:nrow(df),df$depHours)

myProbs <- c(0.50,0.75,0.90)

doMC::registerDoMC(4)

do_fun <- function(){
  foreach(hour = splits, .combine = cbind) %do% {
    quantile(df[hour,"DepDelay"][[1]], myProbs, na.rm=T)}
}
dopar_fun <- function(){
  foreach(hour = splits, .combine = cbind) %dopar% {
    quantile(df[hour,"DepDelay"][[1]], myProbs, na.rm=T)}
}
benchmark <- microbenchmark::microbenchmark(
  do_fun,
  dopar_fun,
  times= 1000000L
)

```

We can see that, on average, a straight %do% runs 38 ns, while %dopar% runs 40 ns. The difference is too less to worry about for this dataset size, to be honest. It might be different for a much larger dataset though.

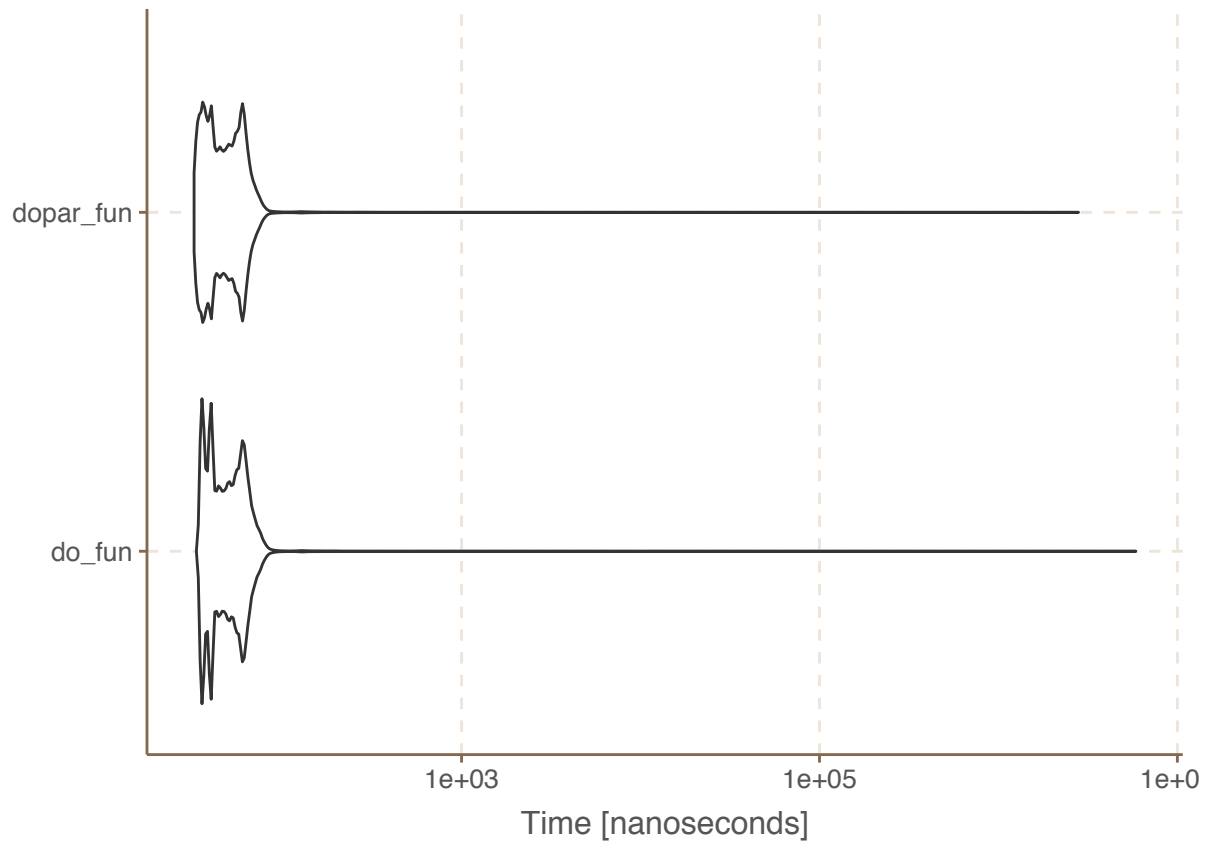
```
benchmark
```

```

## Unit: nanoseconds
##      expr min lq      mean median uq      max neval cld
##   do_fun 33 39 87.37443      48 59 5846084 1e+06   a
## dopar_fun 32 38 72.79741      46 58 2802263 1e+06   a

```

```
microbenchmark::autoplot.microbenchmark(benchmark)
```



## Plots

**Plot #1:** run the analysis as in the book by hour of the day but with `myProbs <- c(0.50,0.75,0.90)`.

Here, the hours are calculated using the `CRSDepTime` variable as shown in the book. Thereafter, for each hour, a `foreach-do` is executed. The plot shows that the delays are the highest for the evening hours, and lowest in the early hours of the morning.


```
df$depHours <- floor(df$CRSDepTime/100)
df$depHours[df$depHours==24] <- 0
splits <- split(1:nrow(df),df$depHours)

myProbs <- c(0.50,0.75,0.90)

delayQuantiles <- foreach(hour = splits, .combine = cbind) %do% {
  quantile(df[hour,"DepDelay"][[1]], myProbs, na.rm=T)
}

colnames(delayQuantiles) <- names(splits)
melted <- reshape2::melt(delayQuantiles)
names(melted) <- c("percentage","hour","delay")
melted %>%
```


```
ggplot(aes(y=delay,x=hour,color=percentage))+
  geom_line()+
  geom_point()
```



**Plot #2:** run the analysis but change it to day of the week and use `myProbs <- c(0.50,0.75,0.90)`.

To get the similar graph for day of the week, we simply need to change the variable used for the splits to `DayOfWeek` as shown below. Here, we can see that the highest delays are 3, 4 and 5 which fall mid-week when the travel might be the highest.

```
splits <- split(1:nrow(df),df$DayOfWeek)
myProbs <- c(0.50,0.75,0.90)
delayQuantiles <- foreach(day = splits, .combine = cbind) %do% {
  quantile(df[day,"DepDelay"][[1]], myProbs, na.rm=T)
}
colnames(delayQuantiles) <- names(splits)
melted <- reshape2::melt(delayQuantiles)
names(melted) <- c("percentage","day","delay")
melted %>%
  ggplot(aes(y=delay,x=day,color=percentage))+
  geom_line()+
  geom_point()
```



**Plot #3:** run the analysis but change it to month of the year and use `myProbs <- c(0.50,0.75,0.90)`.

Now, we can use the Month variable for the splits. The same code as before is used. Now, we see that the delays are highest in Dec and Jan, during the holiday season, when travel is the highest.

```

splits <- split(1:nrow(df), df$Month)
myProbs <- c(0.50, 0.75, 0.90)
delayQuantiles <- foreach(month = splits, .combine = cbind) %do% {
  quantile(df[month, "DepDelay"][[1]], myProbs, na.rm=T)
}
colnames(delayQuantiles) <- names(splits)
melted <- reshape2::melt(delayQuantiles)
names(melted) <- c("percentage", "month", "delay")
melted %>%
  ggplot(aes(y=delay, x=month, color=percentage)) +
  geom_line() +
  geom_point() +
  scale_x_continuous(breaks = seq(1, 12, 1))

```

